



Which model should I use?

How to pick the right AI for the job, and how long to let it think about it.

FREE – PROMPT TEARDOWN FIELD GUIDE

prompt teardown

PROMPTTEARDOWN.COM

PT/01 · FIELD GUIDE

Which model should I use?

*Written for people who want to get more done without
watching their credits vanish by Tuesday afternoon.*

FIRST EDITION · 2026

© PROMPT TEARDOWN · [PROMPTTEARDOWN.COM](https://promptteardown.com)

Contents.

01	Do you actually need to worry about this?	04
	<i>Whether you will ever touch the ceiling</i>	
02	The 2 dials, which model and how long it thinks	06
	<i>The two things you control on every job</i>	
03	The 3 people you call to get the job done	08
	<i>Staff, manager, outside consultant</i>	
04	The effort dial, from a quick glance to a long stare	10
	<i>How long you let the model think</i>	
05	The work, job by job	12
	<i>Errands, writing, decks, spreadsheets, builds</i>	
06	The smartest model is also the wordiest	15
	<i>The paragraph that pays for itself</i>	
07	Every job on one page	17
	<i>The page to keep</i>	

Do you actually need to worry about this?

Every plan has a ceiling. Whether you ever touch it comes down to how you work as much as what plan you're on.

When you are working with AI, there are two meters running at the same time. A short one that refills the same day, and a weekly one that does not. There is also a third limit: how much one conversation can hold. Every company works this way now, and they will change the numbers every few months, so there is no point memorizing them.

What matters is whether you will ever touch that ceiling. And that depends on the kind of work you do, and the plan you are on.

If you mostly write. Copy, emails, articles, bouncing ideas around. Stick with the strong model. Writing burns through the meter slowest, but on a free plan you will still run out on a long day.

If you build things. Decks, websites, spreadsheets, PDFs, code. Whatever plan you are on makes a big difference in how much you get done before you run out.

Your plan

What happens on a heavy week

Free tier

You hit the wall fast. Stay on the cheap model and save the expensive one for the one hard step.

Standard paid tier


You hit it on a busy week. Watch the meter before a big push.

Top-tier plan


You get the most room. You can still burn a week of it in a single day.

Check yours before you assume you have room.

Both numbers are visible in about two clicks. Worth checking before you assume you have room.



CLAUDE CODE — STATUS BAR



CLAUDE APP — PLAN USAGE LIMITS

FIG. 1.0 — WHERE THE TWO CLOCKS LIVE

Every message you send re-reads the whole conversation from the top. A chat that has been running all afternoon costs more per message than a fresh one, and a few PDFs sitting in the thread make it worse. When a chat gets long, start a new one and paste in a short summary of where you got to. That habit buys you more room than upgrading your plan does.

You are not managing a budget. You are managing which jobs deserve the expensive answer.

The people who never seem to run out are not on some better plan. They just know when to switch.

Knowing when to turn the expensive model on, and when to turn it back down, is a skill in itself.

The 2 dials, which model and how long it thinks.

You control two things on every job. Which model you call, and how long you let that model think before it answers.

The first dial is the model. Sonnet, Opus, Fable. Luna, Terra, Sol.

The second dial is how long that model thinks before it answers. Every app names it differently: effort, thinking, reasoning. Same dial.

Turn it down and the answer comes back faster. Turn it up and the model slows down, works the problem through, and checks itself.

More thinking usually means a better answer. But only up to a point.

Chess players have a saying for it.

Think long, think wrong.

Even Anthropic says so.

Anthropic makes this model, and they say the same thing. Push the effort dial all the way up on ordinary work and the model starts overthinking it. Before you reach for a bigger model, try letting the smaller one think a little longer first.¹

The effort dial costs nothing extra to turn. The company selling the pricier model is telling you to try that dial before you pay for a bigger one.

That advice is aimed at building things. If you mostly write, the dial barely changes the result.²

You are picking which person to call, and how long to let them think.

The rest of this guide is that one sentence, applied to the work you actually do.

1 Anthropic, “[When to adjust the effort parameter](#)” and “[Choosing a model.](#)”

2 R2-Write, arxiv.org/pdf/2604.03004.

The 3 people you call to get the job done.

Every company stacks their models the same three ways. Different names, same three steps.

The staff. Fast, cheap, and good at almost everything that fills a normal day. Start here and you will be right most of the time.

The manager. You go to this one when the work actually matters. A serious document, a harder problem, a plan you are going to build on top of.

The outside consultant. Brilliant, billed by the hour, and only in the building about half the week. You bring this one in for the hard part and then you let them go.

THE THREE STEPS • AS OF 2026

Who you call	Claude	ChatGPT	Rough cost
The staff	Sonnet 5	Luna	2 coins
The manager	Opus 5	Terra	5 coins
The consultant	Fable 5	Sol	10 coins

Google and xAI do it the same way.

A landing page built by the consultant might run you \$10 and look gorgeous. The staff builds that same page for \$2 and it comes out plainer.

If the page just takes bookings for a restaurant, plainer is fine, and plainer often converts better. Save the consultant for jobs where the polish is the actual product.

One question, asked before every job.

Almost every model can do almost every job now. Capability stopped being the question a while ago. So stop asking whether the cheap one is able to, and ask this instead.

Can it? Yes. Is it worth it? Probably not.

Run every job through that one and you will get most of these calls right.

The effort dial, from a quick glance to a long stare.

Every model has this dial. You will not touch it often, but knowing it is there already puts you ahead.

Setting	What it is for
LOW	Simple work in volume. Fast and cheap.
MEDIUM	The everyday setting. Good results without the wait.
HIGH	Real thinking. Hard problems, and the final check.
XHIGH	Long jobs with a lot of steps that run for a while.
MAX	The one decision everything else hangs on. Nothing smaller.

Most of your day belongs in low and medium. And the newer the model, the more true that gets. This year's models do more with a quick glance than last year's did with a long stare, so start lower than feels right and climb only when the answer disappoints you.

A bad habit is leaving it on high all the time. You never feel it on any single answer. You feel it on Tuesday.

Switching models mid-chat takes skill.

You do not have to pick one model and live with it for the whole job.

Type `/model` in the middle of a conversation and change it. Nothing restarts, nothing is lost, and the chat keeps everything you already said.

That one move gives you the pattern this entire guide runs on. Plan with the expensive one. Save the plan somewhere. Then drop down and let the cheap one do the work.

By that point the hard decisions are already made. What is left is typing.

Pay for deciding. Save on doing.

Errands, the work where a smart model is wasted.

Some jobs are just fetching and carrying. Every model presses the same button.

The expensive model and the cheap one bring back the same answer, and the cheap one brings it back faster. When you are going back and forth on something, that speed is the whole benefit. Send these to the staff, on low.

- Web searches and lookups
- Pulling from your calendar, your drive, your other tools
- Quick edits, cleanups, and rewrites
- Summarizing something you already have in front of you

Nobody rents a crane to hang a picture.

Writing, decide with the manager and type with the staff.

Writing feels like it deserves the best model. It mostly does not.

A piece is settled before the first sentence gets written. Who it is for. What it has to say. The angle you are taking.

Make those calls with the manager. Then hand it down.

The drafting, the second version, the trim, the cleanup. All of that is typing, and typing is cheap.

Documents work the same way. Plan it with the manager, save the plan, then tell the staff to take the plan and build it.

Deep research is the odd one. It looks like the hardest job on your list, so people reach for the top model on the top setting. But research runs on finding good sources, not on thinking harder about them. A strong model on a low or medium setting is the best value here, and cranking the dial up mostly buys you second-guessing.

A great piece is decided before it is written. Spend where the deciding happens.

Decks, spreadsheets, and builds.

Slide decks. Two jobs in one, what it says and how it looks. AI is good at the first and weak at the second. When researchers scored AI-built decks, design and layout were the worst part of every system they tested.³ So let the staff build every slide, then buy one pass from the consultant to make it look like a person with taste touched it.

Spreadsheets. The exception to most of this guide. One sheet with light math is fine for anyone. Once formulas reach across several sheets, accuracy falls off a cliff. Simple lookups came back right about 89% of the time. Multi-step work across a whole workbook dropped under 20%.⁴ Light work to the staff. Real workbook work is worth the consultant, and worth checking yourself.

Websites and apps. Ask one question first. If this goes wrong, does it cost me an afternoon or a week? An afternoon means stay cheap and redo it. A week means the plan is worth real money. Get the shape right with the consultant, build the routine parts with the staff, then one pass back up top to check the work.

You do not need a consultant-built website to show a restaurant menu. But when the website is the product, buy the good tools.

³ Slide-generation benchmark, arxiv.org/html/2603.07244v1.

⁴ FinSheet-Bench, [arxiv 2603.07316v1](https://arxiv.org/html/2603.07316v1).

The smartest model is also the wordiest.

The better the model gets, the more it talks. You pay for every word of it.

Ask a yes-or-no question without telling it to answer yes or no, and you get four paragraphs. It repeats your question back to you. It tells you what it is about to do before it does it. Then it explains the thing it just did. Ask for a one-page summary and you get three pages, and the extra two are not wrong, they are just not what you asked for.

It also checks its own work without being told to, which is good, and it will quietly do more than you asked, which is not.

This is worst at the top of the stack. Opus 5 is the current example, and it is the chattiest model we have used.

Turning the effort dial down does not fix it. Lower effort means it thinks less about the same long answer. The dial controls thinking, not talking.

You cannot buy a shorter answer. You have to ask for one.

One instruction, given up front, cuts the length of what comes back by about a fifth. That is a fifth of your ceiling, back, on every answer for the rest of the week.

Make it a skill, not a one-time paste.

Paste it at the top of a chat and it works, for a while. Ten turns later the answers creep back to full length. So keep it somewhere you can fire again in two seconds.

STANDING ORDERS – PASTE AS IS

Lead with the answer. Give me the outcome in the first sentence, then the reasoning underneath it, and stop there.

Do not repeat my question back to me. Do not tell me what you are about to do. Do not summarize what you just did.

Keep caveats to one line. Match the length to the ask, and when I ask for something short, short wins over complete.

Do what I asked for and nothing next to it. If you think the ask is wrong, say so in one sentence and then do it anyway.

Do not add a verification pass unless I ask for one. You already check your own work.

If your tool has skills, use one. Claude does. Save the paragraph as a skill, give it a short name, and call it by name every ten turns or whenever the answers start growing.

If it does not, use the settings. ChatGPT has no saved prompt library. Custom instructions are the closest thing, and they apply to every new chat, which suits this one. They still fade over a long conversation.

The trick that works anywhere. Put it in a text expander. TextExpander, Raycast, or the text replacement built into your Mac and phone. Give it a trigger like ;short and re-firing costs six keystrokes in any tool.

One paragraph, written once. A fifth of your ceiling back, every week, forever.

Every job on one page.

Keep this page and you have the guide.

The job	Who you call	How hard they think
Web search and lookups	Staff	LOW
Pulling from your tools	Staff	LOW
Quick edits and summaries	Staff	LOW
Bouncing around ideas	Manager	LOW
Copy and emails	Manager decides, staff drafts	MEDIUM
Documents and PDFs	Manager plans, staff builds	MEDIUM, THEN LOW
Deep research	Manager	LOW TO MEDIUM
Slide decks	Staff builds, consultant polishes	MEDIUM, THEN HIGH
Spreadsheets, one sheet	Staff	MEDIUM
Spreadsheets, many sheets	Consultant	HIGH
Websites and apps	Consultant plans, staff builds	HIGH, THEN LOW

If you remember one line, make it this one.

You are choosing which person to call, and how long to let them think.

That is all of it.

The staff handles your day. The manager makes the calls that matter. The consultant comes in to figure out the hard part, then should go home.

Almost every model can do almost every job now. So the question stopped being what they are capable of a while ago.

Can it? Yes. Is it worth it? Probably not.

EVERY WEEK, ONE PROMPT TAKEN APART

This guide is the long version of what we do every issue.

Prompt Teardown pulls one real prompt apart each week. What it asks for, why it works, and the rebuild you can use the same day. Model-agnostic, free, no filler.

[PROMPTTEARDOWN.COM](https://promptteardown.com)

NEXT — PT/02 · THE TRAVEL PACK